

No AI touches production without a human's informed voice.

An incident commander that diagnoses on its own, then refuses to change anything until a human reads back a code the model was never shown.

Downtime is expensive. Trust is the bottleneck.

PER HOUR OF DOWNTIME

>\$300k

for over 90% of mid-size and large enterprises.
41% put it above \$1M. ITIC 2024

GLOBAL 2000, PER YEAR

\$400B

about \$200M per company, 9% of profits.
Splunk / Oxford Economics 2024

THE BLOCKER

Approval

An AI that can restart production is an audit finding unless you can show who authorized it.

Two contexts. They do not contain the same information.

YOUR SCREEN

HOTEL

ALPHA

Say the action, the service and the code together:
“Roll back auth-service, Hotel Alpha.”

MODEL CONTEXT

Withheld. The agent cannot say this code, cannot guess it, and cannot act without it. The relay checks the readback and runs the change — not the model.

A code alone never authorizes.

The operator must say the action, the service and the code together. Reading two words off a screen is not the same as understanding the change.

BLIND

The code never enters the model's context.

VOICED

A human proves comprehension out loud.

REFUSES

Declines a fix it knows will not hold, and says why.

LOGGED

Recording, turn timeline and postmortem per authorization.

The claim is checkable against a third party.

A checker reads **AssemblyAI's own record** of every session — prompts, tool schemas, calls, results, agent turns — and searches for the code the operator spoke.

SINCE THE FIX

0

occurrences on the model's side. 6 sessions, 100 turns searched.

BEFORE THE FIX

3

a real leak the checker caught: the post-incident summary quoted the spent code.

RUN IT YOURSELF

ayeops-prove
--selftest

The checker fails on a planted leak and passes a clean session.

Real containers, real crash loops, real recovery.

METRIC	RESULT	METHOD
Bad deploy → all green	29 s	detection to healthy, wall clock
Turn latency	0.84–1.09 s	from the operator’s last voiced frame
Recovery from a cut link	1.8 s	context intact
Scenario tests	39 / 199 assertions	green in CI on every push

Sold as the control that makes an agent approvable.

BUYER

Platform and SRE teams already running agents against infrastructure, in regulated industries where every production change needs an auditable approver.

MODEL

Per-seat for on-call engineers, plus a fee per authorization. The audit bundle is the product: who authorized what, when, and how.

WEDGE

Autonomous remediation is blocked at security review, not at capability. We sell the thing that gets it past review.

Nine voice incident commanders. All compete on speed.

“Destructive mutations require explicit TTL-bounded, single-use tokens.
Spoken keywords never authorize execution.”

– the strongest competing entry, from its own deck

- It removed the language model, and its Kubernetes cluster is an **in-memory simulator**.
- We took the opposite position: **a human’s voice is the only thing that authorizes**.
- And unlike a speed claim, ours is **verifiable by a third party**, not measured by us.

The cluster adapter is 127 lines. Kubernetes is next.

NEXT

Kubernetes and Argo adapters — rollback maps to **rollout undo**, restart to **rollout restart**, scale to **scale**.

THEN

SSO, audit export to the change-management system, and an on-prem relay so the recording never leaves the customer.

TODAY

Live, open source, and checkable.
ayeops.vexralabs.com